Artificial Intelligence and Retrieval-Augmented Generation

Retrieval-Augmented Generation (RAG) combines dense vector retrieval with LLMs.

By anchoring generative responses to verified knowledge bases or PDF documents,

RAG architectures mitigate hallucination and provide deterministic citations.

Key Components:

1. Document Parsing & Chunking (e.g. RecursiveCharacterTextSplitter)
2. Embedding Generation (e.g. all-MiniLM-L6-v2)
3. Vector Store Retrieval (e.g. Qdrant / FAISS)
4. Grounded Synthesis via Large Language Models